

Algorithm	Best Fit Situation	Advantage	Disadvantage
Random Forest	<ol style="list-style-type: none">1. Suited for at almost any machine learning problem2. Ex. Bioinformatics	<ol style="list-style-type: none">1. Can work in parallel2. Susceptible to overfits3. Automatically handles missing values if you impute using a special number4. No need to transform any variable5. No need to tweak parameters	<ol style="list-style-type: none">6. Difficult to interpret7. Weaker on regression when estimating values at the extremities of the distribution of response values8. Biased in multiclass problems toward more frequent classes
Gradient Boosting	<ol style="list-style-type: none">1. Apt at almost any machine learning problem2. Search engines (solving the problem of learning to rank)	<ol style="list-style-type: none">3. It can approximate most nonlinear function4. Best in class predictor5. Automatically handles missing values6. No need to transform any variable	<ol style="list-style-type: none">7. It can overfit if run for too many iterations8. Sensitive to noisy data and outliers9. Doesn't work at its best without parameter tuning
Linear Regression	<ol style="list-style-type: none">1. Baseline prediction2. Econometric predictions	<ol style="list-style-type: none">1. Simple to understand and explain	<ol style="list-style-type: none">1. You have to work hard to make it fit

	<ol style="list-style-type: none">3. Modelling4. marketing responses	<p>It seldom overfits</p> <ol style="list-style-type: none">2. Using L1 & L2 regularization is effective in feature selection3. Fast to train4. Easy to train on big data thanks to its stochastic version	<p>nonlinear functions</p> <ol style="list-style-type: none">2. Can suffer from outliers
Support Vector Machine	<ol style="list-style-type: none">1. Character recognition2. Image recognition3. Text classification	<ol style="list-style-type: none">1. Automatic non-linear feature creation2. Can approximate complex non-linear functions3. Works only with a portion of the examples (the support vectors)	<ol style="list-style-type: none">1. Difficult to interpret when applying non-linear kernels2. Suffers from too many examples, after 10,000 examples it starts taking too long to train
K Nearest Neighbor	<ol style="list-style-type: none">1. Computer vision2. Multilabel tagging3. Recommender systems4. Spell checking problems	<ol style="list-style-type: none">1. Fast, lazy training2. Can naturally handle extreme multiclass problems (like tagging text)	<ol style="list-style-type: none">1. Slow and cumbersome in the predicting phase2. Can fail to predict correctly due to the curse of dimensionality

Adaboost	Facedetection	<ol style="list-style-type: none">1. Automatically handles missing values2. No need to transform any variable3. It doesn't overfit easily4. Few parameters to tweak5. It can leverage many different weak-learners	<ol style="list-style-type: none">1. Sensitive to noisy data and outliers2. Never the best in class predictions
Naive Bayes	<ol style="list-style-type: none">1. Face recognition2. Sentiment analysis3. Spam detection4. Text classification	<ol style="list-style-type: none">1. Easy and fast to implement, doesn't require too much memory and can be used for online learning2. Easy to understand3. Takes into account prior knowledge	<ol style="list-style-type: none">1. Strong and unrealistic feature independence assumptions2. Fails estimating rare occurrences3. Suffers from irrelevant features
Neural Networks	<ol style="list-style-type: none">1. Image recognition2. Language recognition and translation3. Speech recognition	<ol style="list-style-type: none">1. It can approximate any non-linear function	<ol style="list-style-type: none">1. It requires you to define a network architecture

	<ol style="list-style-type: none">4. Vision recognition	<ol style="list-style-type: none">2. Robust to outliers3. It can work with image, text and sound data	<ol style="list-style-type: none">2. Difficult to tune because of too many parameters and you have also to decide the architecture of the network3. Difficult to interpret4. Easy to overfit
Logistic Regression	<ol style="list-style-type: none">1. Ordering results by probability2. Modelling marketing responses	<ol style="list-style-type: none">1. Simple to understand and explain2. It seldom overfits3. Using L1 & L2 regularization is effective in feature selection4. The best algorithm for predicting probabilities of an event5. Fast to train6. Easy to train on big data thanks to its stochastic version	<ol style="list-style-type: none">1. You have to work hard to make it fit non-linear functions2. Can suffer from outliers

SVD	Recommendation System	1. Can restructure data in a meaningful way	1. Difficult to understand why data has been restructured in a certain way
PCA	1. Removing collinearity 2. Reducing dimensions of the dataset	1. Can reduce data dimensionality	1. Implies strong linear assumptions (components are a weighted summations of features)
K means	1. Segmentation	1. Fast in finding clusters 2. Can detect outliers in multiple dimensions	1. Suffers from multicollinearit 2. Clusters are spherical, can't detect groups of other shape 3. Unstable solutions, depends on initialization